

## EXPERIENCE

### UCLA NLP

*June 2019 - Present*

*ML/NLP Researcher*

- Used BERT and unsupervised contrastive learning to create sentence embeddings.
- Investigated what causes catastrophic forgetting when BERT fine-tuning.
- Supervised by Professor Kai-Wei Chang.

### Blend

*September 2017 - September 2018*

*Machine Learning Engineer*

- Deployed a LSTM-RNN model to identify app glitches in user comments, improved P/R and accuracy from 70% to 90% from previous keyword based method.
- Applied deep learning NLP methods to categorize and cluster user reviews.
- Created and maintained lambda architecture event/log ingestion pipeline with PySpark.

### Blend

*June 2017 - September 2017*

*Software Engineering Intern*

- Trained an RNN to [predict loan application submission](#) with 90% accuracy.
- Deployed a multi-user data science toolbox with PySpark support.
- Deployed Airflow with CI/CD to standardize analytics and coordinate ETL jobs.

### NASA JPL

*June 2016 - January 2017*

*Software Engineering Intern*

- Used DeepDive and MITIE to create RDF triples for a natural language QA system.
- Created several APIs/backends with Flask, Docker, and Elasticsearch.

---

## EDUCATION

### UCLA

*Expected December 2019*

*B.S. Computer Science, Upper Division GPA: 3.57*

- Coursework: ECE 239 Deep Learning, CS 269 Natural Language Processing, CS 267A Probabilistic Programming & Relational Learning
- Technical Breadth: Mathematics (Linear Algebra, Real Analysis, Optimization)
- ACM AI President (2017). Received ACM Student Chapter Excellence Award.

---

## PROJECTS

### [QuickThoughts](#)

- PyTorch reimplementation of QuickThoughts paper by Logeswaran and Lee.
- Achieves SOTA performance on several downstream classification tasks.
- Theoretical analysis on the effect of context size on performance [here](#).

### [DistBelief](#)

- Implemented DownpourSGD (asynchronous distributed training) in PyTorch.
- Reduced training time for AlexNet on CIFAR10 by 50 minutes (45%) vs. single-node SGD
- Wrote small message passing framework on top of PyTorch's native distributed point to point communication.

### [PDB Evaluator](#)

- Wrote query parser and first-order lifted inference for probabilistic database queries.
- Used SQLite to achieve up to 4x speedup on inference on decomposable queries.